

Cuckoo Hashing using CUDA

Hongtu Xu, 2019533229

xuht1@shanghaitech.edu.cn

January 1, 2022

Abstract

In computing, hash table is a widely used data structure that implements an associative array abstract structure, a structure that maps keys to values. Cuckoo Hashing is one among the hashing techniques which provides high memory usage with constant time access [1]. At the same time, GPU makes high throughput hash table possible. In this lab, I implement Cuckoo Hashing use CUDA as well as testing correctness and doing experiments to analyze performance.

1 Introduction

1.1 Compile & Build

You need a C++ compiler which supports at least C++ 17 standard and nvcc which supports at least C++ 17 standard for CUDA C++.

I use CMake to build the project, so you also need to have CMake with version at least 3.21 to compile our codes.

Compile by:

```
mkdir build
cd .\build
cmake ..
cmake --build . --config Release
```

Run the benchmark by main target under tests:

```
cd .\build\tests
.\main.exe
```

It will automatically run the four experiments with an additional correctness test.

1.2 Interface

We provide a function `insert_and_lookup` to use our hash table.

It will first insert the given keys into the hash table and then do the lookup and store the result according to the given lookup key sequence.

1.3 Benchmark Configures

I run my code on my own computer. The detailed information is:

| Config | Value |
|----------------|----------------------------------|
| CPU | AMD Ryzen 7 3700X |
| GPU | NVIDIA GeForce RTX 2060 SUPER |
| Memory | 32 GB DDR4 3000 MHz 2 of 4 slots |
| GPU Memory | 8 GB GDDR6 |
| OS | Windows 11 Pro 22000.37 |
| Compiler | MSVC 19.30.30706, nvcc V11.5.119 |
| Optimize | MSVC: /O2, nvcc: -O3 |
| CUDA Cores | 2176 |
| CUDA Version | 11.5 |
| Driver Version | GeForce Game Ready Driver 497.29 |

To test scalability, I also run my code on a cloud virtual machine provided by tencent cloud. The detailed information is:

| Config | Value |
|----------------|-------------------------------|
| CPU | AMD EPYC ROME™ 7k62 (16 vCPU) |
| GPU | NVIDIA TESLA A100 40GB PCIe |
| Memory | 96 GB |
| GPU Memory | 40 GB HBM2e |
| OS | Ubuntu 20.04 LTS |
| Compiler | gcc 11.1.0, nvcc V11.1.105 |
| Optimize | gcc: -O2, nvcc: -O3 |
| CUDA Cores | 6912 |
| CUDA Version | 11.1 |
| Driver Version | 450.102.04 |

2 Implement Details

2.1 Cuckoo Hashing Algorithm

Cuckoo Hashing is one of the hash table schemas which provides high memory utilization and constant access time [1]. It is named after a type of bird who lays its eggs in other birds' nests, and

whose chicks push out the other eggs from the nest after they hatch.

Unlike many other hash tables, Cuckoo Hashing uses t hash functions where $t > 1$. Cuckoo hashing maintains t tables, each key k can be stored into t locations corresponds to the location that uses hash function h_i . In the insertion process, suppose we want to insert a key k into the hash table and we have t hash functions h_0, h_1, \dots, h_{t-1} while the tables are T_0, T_1, \dots, T_{t-1} . We will first use hash function h_1 to try to insert k into $T_1[h_1(k)]$. If the cell is already occupied, then we will evict the key in that slot k' , put the k into that cell and try to insert k' into tables after circularly. If we still need to evict key after a bound times, then we need to choose some other functions and rehash.

Algorithm 1 Cuckoo Hashing Insert

h_1, h_2, \dots, h_t are t hash functions
 T_1, T_2, \dots, T_t are corresponding tables
 C is the size of the table (capacity)
input: a key k
for $i = 0 \rightarrow \text{iter-bound}$ **do**
 $\text{loc} \leftarrow h_i \bmod t(k) \bmod C$
 if $T_i \bmod t[\text{loc}] = \text{empty}$ **then**
 $T_i \bmod t[\text{loc}] \leftarrow k$
 return
 else
 $\text{swap}(k, T_i \bmod t[\text{loc}])$
 end if
end for
 $\text{rehash}()$
 $\text{insert}(k)$

The lookup process is much simpler, we just need to loop through all t possible locations and compare the key with the stored key in the table whether they are equal. Obviously, this process is done in $O(t)$ and since t is a constant, then Cuckoo Hashing provides constant time access.

Algorithm 2 Cuckoo Hashing Lookup

h_1, h_2, \dots, h_t are t hash functions
 T_1, T_2, \dots, T_t are corresponding tables
 C is the size of the table (capacity)
input: a key k
for $i = 0 \rightarrow t - 1$ **do**
 $\text{loc} = h_i(k) \bmod C$
 if $T_i[\text{loc}] = k$ **then**
 return true
 end if
end for
return false

We can find that the lookup time in Cuckoo Hashing is $O(t) = O(1)$, which is very good. But

the worst case exists when key is not inserted into the hash table, it will loop through all t hash functions.

2.2 Table Storage

Since Cuckoo Hashing needs to maintain t tables where t is the number of hash functions. From the algorithm, it is obvious that the tables are independent to each other. So, instead of storing all the t tables into a big linear array, we choose to store t linear arrays, which performs better because of good pipeline and cache utilization.

2.3 Memory Management

For hash table, we mainly need some big linear arrays. Although unified memory can greatly decrease the difficulty for memory management between device and host. But it needs some techniques to achieve good performance and it may still perform a little bit worse than directly manage host and device memory. Therefore, I create a class to encapsulate cuda array. By overloading assignment operator, I can make copy operation between device and host simpler. Also, by encapsulation, I can allocate and free the memory by some class methods easily.

2.4 Random Set Generation

We need to generate random set to test our hash table. Since it is a set which means it cannot contain the same value, then I choose to first generate a such set on CPU and then copy it to GPU. I use `std::mt19937`, a mersenn twister engine, as our random number engine. To make the generation faster, I create a bitset whose size is 2^{32} , which will cost 0.5 GB memory but provides much better performance than using a hash table to filter out duplicates.

For the lookup set generation, I first generate a totally different set respected to the original random set with the same size. Then, I implement a random shuffle kernel function to support shuffle operation on a device array. For the lookup set, I will first shuffle it randomly then overwrite some keys using elements from the different set according to the given percentage. This is much faster than using bitset or hash table to filter and create another set.

2.5 Multi-level Table

One possible method to improve the performance of our GPU hash table is to use multi-level table [2]. The bottleneck of Cuckoo Hashing is memory access, uncoalesced global memory access will greatly influence our perfor-

mance. So, we may create multi-level tables. For example, two-level Cuckoo Hashing, we first hash the elements into different buckets and then perform Cuckoo Hashing locally inside those buckets.

This often performs better than single-level table when rehashing is needed. However, in our experiments, we do not need to rehash that much and according to the result in [2], the multi-level table performs even worse than single-level. Therefore, I choose to not use multi-level table and in my experiments, I compared my code with an existing multi-level cuckoo hashing code and I performs much better than him.

2.6 Parallel d-Pipeline

During researching, I found another way to increase the throughput, which is parallel d-pipeline [3]. Consider the insertion process of one key k , it will go through tables like a pipeline:

input $\rightarrow T_0 \rightarrow T_1 \rightarrow T_2 \rightarrow \dots \rightarrow T_{t-1} \rightarrow \text{rehash}$

If success at one table in the middle, it will jump out the pipeline. Therefore, when a key failed at T_0 and evict one to try to insert in T_1 , another new key can start to insert at T_0 without any race condition. And this forms a pipeline.

However, implement this directly is hard in coding. Hopefully, by taking advantage of modern compiler and C++ we can form the pipeline using template metaprogramming to expand the code and finally form a good pipeline through compiler optimization.

2.7 Hash Function

Since Cuckoo Hashing has a constant lookup time, then we just need to focus on insertion. A good hash function can cause less collisions and thus we can perform fewer evictions and finally gain a better performance in insertion.

2.7.1 Universal Hashing

Universal Hashing takes advantages of randomness. No matter what n input keys are, every operation takes $O(n/m)$ time in expectation, for a size m hash table, which is optimal if no randomness. Instead of using a fixed hash functions, universal hashing uses a random hash function, chosen from a universal hash family H that satisfied:

$$\mathbb{P}_{h \in H}(h(x) = h(y)) = \frac{1}{m}$$

We can construct hash function h_{ab} by choosing a prime number p s.t. $p > m$ and $p >$ all keys:

$$h_{ab}(k) = ((ak + b) \bmod p) \bmod m$$

Let $H_{pm} = \{h_{ab} | a \in \{1, 2, \dots, p-1\}, b \in \{0, 1, \dots, p-1\}\}$, then H_{pm} is a universal hash family. We can randomly choose h_{ab} from H_{pm} . However, we need to ensure $p > m$ and $p >$ all keys. In our experiments, the maximum value of key is the maximum value of 32-bit integer, which means we need to choose a prime $> 2^{32}$ and store it in 64-bit integer. 64-bit integer calculations are too slow in GPU, and if we use smaller p within 32-bit integer, the performance is not good. Therefore, I do not use this hash function at last.

2.7.2 Perfect Hashing

This is the ideal hash function because it maps different key to exactly different values. However, minimal perfect hashing costs too much time during construction [1] while some other variants of perfect hashing like achieving this by take two level universal hashing, are hard to use in Cuckoo Hashing and cannot support newly come keys.

2.7.3 xxHash

xxHash is an extremely fast Hash algorithm, running at RAM speed limits [4]. Besides extremely high performance, it provides excellent quality which passes Google hash tests. After testing different hash functions, xxHash performs better in most cases. So, I use xxHash as the hash functions at last.

2.8 Template metaprogramming

Template metaprogramming is good way to transform some dynamic process into static process as well as gaining further optimization by compiler and thus have a better performance.

We pass the seed of xxHash through template, by doing this, all the multiplication involving the seed will be determined at compile time and after compiler optimization the code will become extremely small and simple and will be inlined. Similarly, we create a template for insert and lookup, the recursion in template will be expand at compile time and form a pipeline, we can do this because the maximum eviction bound can be chosen as a constant.

By taking advantage of modern compiler, we gain a much higher performance with less cost at coding level.

2.9 Timer

To measure the time precisely, I use `cudaEvent` to measure the time. I record two `cudaEvent` before and after invoking the kernel function, respectively. Before calculating the time elapsed, we need to first synchronize the event to ensure the kernel function has already finished. Then use `cudaEventTimeElapsed` to calculate the time elapsed. To obtain reliable results, I perform each experiment 5 times.

2.10 CUDA Block Size Tuning

I tuned the CUDA block size with NVIDIA nsight compute. From Figure 2.10, we can see that we reach highest occupancy with block size of 512 or 1024. After further benchmark test, I finally choose 512 as the block size, which obtains a little bit better performance than choosing 1024. Because we may have more registers per thread to use when choosing 512 as the block size.


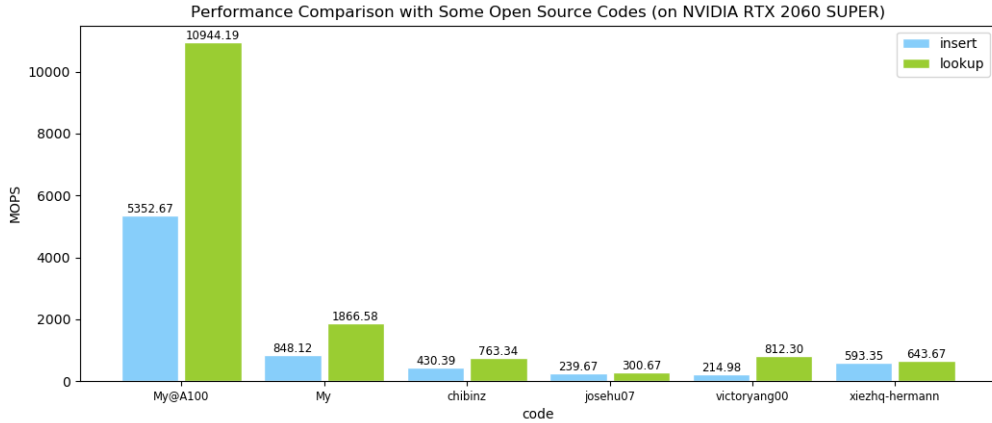


Figure 2.10: Impact of varying block size to warp occupancy

3 Results

After running experiments, the benchmark result is listed as follows:



Performance comparison with some open source codes

3.1 Experiment 1: Insertion

In this experiment, I create a hash table of size 2^{25} in GPU global memory, where each table entry stores a 32-bit integer. I insert a set of 2^s random integer keys into the hash table, for $s = 10, 11, \dots, 24$. And I test for both using $t = 2$ hash functions and $t = 3$ hash functions. To produce reliable result, each experiment are repeated 5 times.

We can find that using 3 hash functions performs better than using 2 hash functions in most cases because fewer evictions are needed when

using 3 hash functions. 3 hash function performs especially better when s is larger, when $s = 24$ the difference is extremely obvious.

With NVIDIA TESLA A100, we get our peak insertion performance of **5352.674049** millions of insertions per second. The result between different number of hash functions stays the same, that is $t = 3$ performs better than $t = 2$ also due to the decrease in the number of evictions.

We generally get more than $6\times$ speedup on TESLA A100 with respected to RTX 2060 SUPER, which shows that our implementation ob-

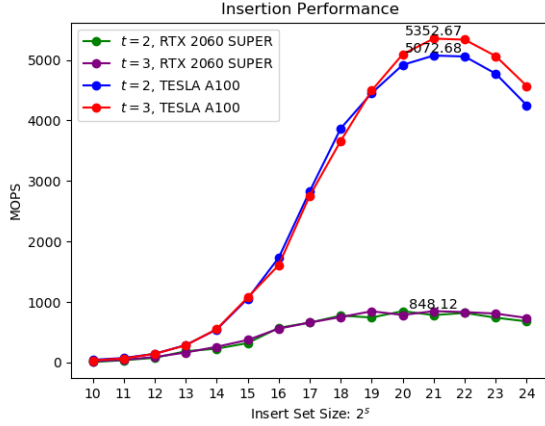


Figure 1: Experiment 1 Performance

| s | t | MOPS | Time (ms) | StdDev |
|-----|-----|-------------------|-----------|----------|
| 10 | 2 | 7.446709 | 0.137510 | 0.063818 |
| 11 | 2 | 37.200651 | 0.055053 | 0.027095 |
| 12 | 2 | 76.472696 | 0.053562 | 0.016374 |
| 13 | 2 | 185.561035 | 0.044147 | 0.015856 |
| 14 | 2 | 227.859371 | 0.071904 | 0.013538 |
| 15 | 2 | 320.661368 | 0.102189 | 0.049058 |
| 16 | 2 | 572.194902 | 0.114534 | 0.009390 |
| 17 | 2 | 658.499736 | 0.199046 | 0.009860 |
| 18 | 2 | 773.983850 | 0.338694 | 0.005635 |
| 19 | 2 | 743.422896 | 0.705235 | 0.137846 |
| 20 | 2 | 848.028746 | 1.236486 | 0.004866 |
| 21 | 2 | 782.824122 | 2.678957 | 0.274631 |
| 22 | 2 | 818.832553 | 5.122298 | 0.240417 |
| 23 | 2 | 740.279704 | 11.331674 | 0.807136 |
| 24 | 2 | 678.550312 | 24.725088 | 0.903823 |

Table 1: Insertion test on NVIDIA GeForce RTX 2060 SUPER using $t = 2$ hash functions

tain an excellent scalability because we mainly benefit from the GPU memory bandwidth and the HBM2e is 3.47 times as fast as GDDR6 while A100 takes 2.7 times higher in TFLOPS.

3.2 Experiment 2: Lookup

In this experiment, I insert a set S of 2^{24} random keys into a hash table of size 2^{25} , then perform lookups for the following sets of keys S_0, \dots, S_{10} . Each set S_i contains 2^{24} keys, where $(100 - 10i)$ percent of the keys are randomly chosen from S , and the remainder are random 32-bit keys. To produce reliable result, each experiment are repeated 5 times.

From table 5 and 6 we can find that for the lookup test, time increases as i increases because more elements are not in the hash table as i becomes bigger. The lookup time in Cuckoo Hashing is $O(t)$ when the element we try to lookup is

| s | t | MOPS | Time (ms) | StdDev |
|-----|-----|-------------------|-----------|----------|
| 10 | 3 | 21.214532 | 0.048269 | 0.009736 |
| 11 | 3 | 44.071065 | 0.046470 | 0.002681 |
| 12 | 3 | 90.026727 | 0.045498 | 0.004228 |
| 13 | 3 | 164.018451 | 0.049946 | 0.009279 |
| 14 | 3 | 258.638111 | 0.063347 | 0.005503 |
| 15 | 3 | 371.984887 | 0.088090 | 0.015388 |
| 16 | 3 | 559.348878 | 0.117165 | 0.003982 |
| 17 | 3 | 659.666305 | 0.198694 | 0.017240 |
| 18 | 3 | 750.733137 | 0.349184 | 0.026919 |
| 19 | 3 | 846.473374 | 0.619379 | 0.001562 |
| 20 | 3 | 784.772035 | 1.336154 | 0.234942 |
| 21 | 3 | 848.118706 | 2.472710 | 0.099679 |
| 22 | 3 | 832.881107 | 5.035898 | 0.234401 |
| 23 | 3 | 807.912957 | 10.383059 | 0.109445 |
| 24 | 3 | 736.136639 | 22.790899 | 0.155948 |

Table 2: Insertion test on NVIDIA GeForce RTX 2060 SUPER using 3 hash functions

| s | t | MOPS | Time (ms) | StdDev |
|-----|-----|--------------------|-----------|----------|
| 10 | 2 | 44.211108 | 0.023162 | 0.003254 |
| 11 | 2 | 73.613987 | 0.027821 | 0.002437 |
| 12 | 2 | 143.465591 | 0.028550 | 0.002464 |
| 13 | 2 | 286.802600 | 0.028563 | 0.002506 |
| 14 | 2 | 545.028746 | 0.030061 | 0.000516 |
| 15 | 2 | 1055.452489 | 0.031046 | 0.000901 |
| 16 | 2 | 1730.606726 | 0.037869 | 0.002185 |
| 17 | 2 | 2829.120113 | 0.046330 | 0.001364 |
| 18 | 2 | 3867.799868 | 0.067776 | 0.003297 |
| 19 | 2 | 4452.415964 | 0.117754 | 0.001905 |
| 20 | 2 | 4918.495447 | 0.213190 | 0.001987 |
| 21 | 2 | 5072.681438 | 0.413421 | 0.001089 |
| 22 | 2 | 5057.063274 | 0.829395 | 0.002853 |
| 23 | 2 | 4773.093010 | 1.757478 | 0.002250 |
| 24 | 2 | 4252.760722 | 3.945018 | 0.004734 |

Table 3: Insertion test on NVIDIA TESLA A100 40GB PCIe using $t = 2$ hash functions

already in the hash table, then it will take less than t steps to find that key. Otherwise, it will take the whole t step to give a not found.

Therefore, the performance of lookup reaches the peak when all the lookup keys are in the hash table. While all the keys are not in the hash table, it takes the longest time but the time complexity is still a constant because the number of hash functions t is a constant. And when t is smaller, the performance is better. However, smaller t may increase the number of evictions in insertion. So, in practice we need balance between insertion and lookup.

According to Table 6, we get a peak performance of **10944.190395** millions of lookups per second with NVIDIA TESLA A100, which is

| s | t | MOPS | Time (ms) | StdDev |
|-----|-----|--------------------|-----------|----------|
| 10 | 3 | 34.453057 | 0.029722 | 0.001653 |
| 11 | 3 | 67.100022 | 0.030522 | 0.001605 |
| 12 | 3 | 142.793397 | 0.028685 | 0.003427 |
| 13 | 3 | 283.248506 | 0.028922 | 0.000490 |
| 14 | 3 | 551.724145 | 0.029696 | 0.001909 |
| 15 | 3 | 1073.600350 | 0.030522 | 0.001782 |
| 16 | 3 | 1609.303782 | 0.040723 | 0.001657 |
| 17 | 3 | 2755.651233 | 0.047565 | 0.002297 |
| 18 | 3 | 3661.065435 | 0.071603 | 0.003161 |
| 19 | 3 | 4495.664528 | 0.116621 | 0.004782 |
| 20 | 3 | 5094.368927 | 0.205830 | 0.003153 |
| 21 | 3 | 5352.674049 | 0.391795 | 0.003084 |
| 22 | 3 | 5335.417389 | 0.786125 | 0.001248 |
| 23 | 3 | 5063.862307 | 1.656563 | 0.004728 |
| 24 | 3 | 4570.751084 | 3.670560 | 0.005147 |

Table 4: Insertion test on NVIDIA TESLA A100 40GB PCIe using $t = 3$ hash functions

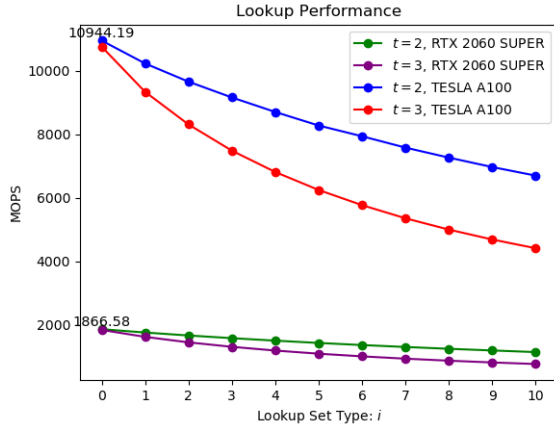


Figure 2: Experiment 2 Performance

extremely fast while in the worst case all the lookups are done in 2.51ms when $t = 2$. In addition, our implementation still obtains an excellent in the lookup, reaching a speedup of more $6\times$ with respected to RTX 2060 SUPER.

3.3 Experiment 3: Size

In this experiment, I create a fixed set of $n = 2^{24}$ random keys and measure the time to insert the keys into hash tables of size sn where $s = 1.01, 1.02, 1.05, 1.1, 1.2, \dots, 2$. I terminate the experiment if the insertion takes too long which means the insert process of Cuckoo Hashing is hard to complete. To produce reliable result, each experiment are repeated 5 times.

According to Figure 3, we can see that generally, the performance increases as the size increases because there will be less evictions and rehashes.

| i | t | MOPS | Time (ms) | StdDev |
|-----|-----|--------------------|-----------|----------|
| 0 | 2 | 1866.584072 | 8.988192 | 0.107335 |
| 1 | 2 | 1761.835282 | 9.522579 | 0.007990 |
| 2 | 2 | 1668.597444 | 10.054682 | 0.004976 |
| 3 | 2 | 1584.623803 | 10.587507 | 0.003354 |
| 4 | 2 | 1506.272597 | 11.138234 | 0.001299 |
| 5 | 2 | 1433.399665 | 11.704493 | 0.008614 |
| 6 | 2 | 1367.833913 | 12.265536 | 0.005892 |
| 7 | 2 | 1306.748787 | 12.838899 | 0.005308 |
| 8 | 2 | 1249.908829 | 13.422752 | 0.004688 |
| 9 | 2 | 1197.400414 | 14.011366 | 0.002645 |
| 10 | 2 | 1146.550585 | 14.632774 | 0.057043 |
| 0 | 3 | 1843.989880 | 9.098323 | 0.002369 |
| 1 | 3 | 1625.353454 | 10.322195 | 0.002685 |
| 2 | 3 | 1452.115347 | 11.553639 | 0.010561 |
| 3 | 3 | 1311.857373 | 12.788902 | 0.005694 |
| 4 | 3 | 1194.372580 | 14.046886 | 0.010308 |
| 5 | 3 | 1095.647198 | 15.312608 | 0.003277 |
| 6 | 3 | 1011.771790 | 16.582016 | 0.004829 |
| 7 | 3 | 939.055437 | 17.866055 | 0.005416 |
| 8 | 3 | 875.845898 | 19.155443 | 0.005154 |
| 9 | 3 | 820.318196 | 20.452083 | 0.006410 |
| 10 | 3 | 771.261283 | 21.752960 | 0.007894 |

Table 5: Lookup test on NVIDIA GeForce RTX 2060 SUPER using 2 and 3 hash functions

| i | t | MOPS | Time (ms) | StdDev |
|-----|-----|---------------------|-----------|----------|
| 0 | 2 | 10944.190395 | 1.532979 | 0.001634 |
| 1 | 2 | 10225.301127 | 1.640755 | 0.033481 |
| 2 | 2 | 9652.516268 | 1.738118 | 0.039536 |
| 3 | 2 | 9155.470175 | 1.832480 | 0.039523 |
| 4 | 2 | 8699.132169 | 1.928608 | 0.042077 |
| 5 | 2 | 8272.762610 | 2.028006 | 0.048818 |
| 6 | 2 | 7935.005113 | 2.114330 | 0.025525 |
| 7 | 2 | 7577.905498 | 2.213965 | 0.029892 |
| 8 | 2 | 7263.175951 | 2.309901 | 0.026111 |
| 9 | 2 | 6966.356637 | 2.408320 | 0.027162 |
| 10 | 2 | 6697.239074 | 2.505094 | 0.027777 |
| 0 | 3 | 10752.596504 | 1.560294 | 0.019172 |
| 1 | 3 | 9331.956724 | 1.797824 | 0.045214 |
| 2 | 3 | 8307.342176 | 2.019565 | 0.055082 |
| 3 | 3 | 7482.538448 | 2.242182 | 0.065040 |
| 4 | 3 | 6810.226405 | 2.463533 | 0.072663 |
| 5 | 3 | 6244.869361 | 2.686560 | 0.083055 |
| 6 | 3 | 5768.100057 | 2.908621 | 0.093335 |
| 7 | 3 | 5354.653074 | 3.133203 | 0.101179 |
| 8 | 3 | 4998.941599 | 3.356154 | 0.112646 |
| 9 | 3 | 4687.387080 | 3.579226 | 0.121195 |
| 10 | 3 | 4413.198677 | 3.801600 | 0.130535 |

Table 6: Lookup test on NVIDIA TESLA A100 using 2 and 3 hash functions

When size is $1.7n$, we get a performance drop because we encounter one rehash in this test

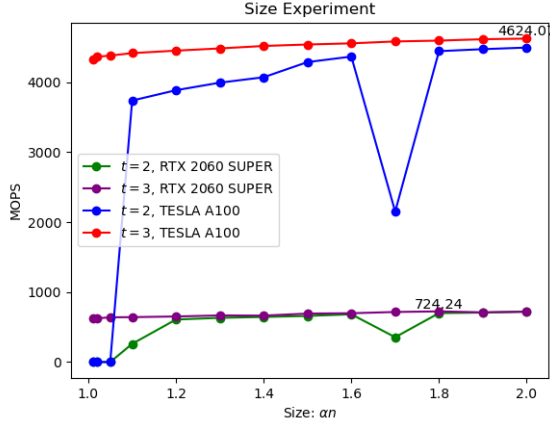


Figure 3: Experiment 3: Size Test

| s | t | MOPS | Time (ms) | StdDev |
|------|-----|-------------------|-----------|----------|
| 2.0 | 2 | 718.569511 | 23.348077 | 0.526595 |
| 1.9 | 2 | 707.180415 | 23.724096 | 0.583006 |
| 1.8 | 2 | 699.033878 | 24.000576 | 0.654333 |
| 1.7 | 2 | 354.465352 | 47.331046 | 1.299906 |
| 1.6 | 2 | 682.258901 | 24.590688 | 0.316653 |
| 1.5 | 2 | 658.699933 | 25.470195 | 0.726411 |
| 1.4 | 2 | 644.583205 | 26.028007 | 0.887359 |
| 1.3 | 2 | 631.628612 | 26.561837 | 0.594943 |
| 1.2 | 2 | 607.410984 | 27.620864 | 0.922161 |
| 1.1 | 2 | 259.893430 | 64.554214 | 1.874496 |
| 1.05 | 2 | N/A | N/A | N/A |
| 1.02 | 2 | N/A | N/A | N/A |
| 1.01 | 2 | N/A | N/A | N/A |

Table 7: Size test on NVIDIA GeForce RTX 2060 SUPER using 2 hash functions

key set. When the size $\leq 1.05n$ and if we choose 2 hash functions, the hash table will infinitely rehash.

When we choose 3 hash functions, there will not be any rehash and the performance increases as the size increases due to fewer evictions.

According to Figure 3 and Table 10, on NVIDIA TESLA A100, we get our peak performance of **4624.068615** with 3 hash functions and size of $2.0n$.

3.4 Experiment 4: Bound

In this experiment, I use $n = 2^{24}$ random keys and a hash table of size $1.4n$. I test Cuckoo Hashing with different bounds of $l \log n$ on the maximum length of eviction chain before restarting.

I terminate the experiment if the insertion takes too long which means the insert process of Cuckoo Hashing is hard to complete. To pro-

| s | t | MOPS | Time (ms) | StdDev |
|------|-----|-------------------|-----------|----------|
| 2.0 | 3 | 720.760163 | 23.277113 | 0.823124 |
| 1.9 | 3 | 710.266831 | 23.621004 | 0.795786 |
| 1.8 | 3 | 724.240721 | 23.165248 | 0.162821 |
| 1.7 | 3 | 715.366934 | 23.452602 | 0.253312 |
| 1.6 | 3 | 697.417215 | 24.056211 | 0.665444 |
| 1.5 | 3 | 692.596413 | 24.223654 | 0.630849 |
| 1.4 | 3 | 663.355432 | 25.291443 | 0.560652 |
| 1.3 | 3 | 665.923213 | 25.193920 | 0.797485 |
| 1.2 | 3 | 651.055876 | 25.769241 | 0.822991 |
| 1.1 | 3 | 641.431637 | 26.155891 | 0.766913 |
| 1.05 | 3 | 639.108578 | 26.250964 | 0.603821 |
| 1.02 | 3 | 624.410461 | 26.868890 | 0.688250 |
| 1.01 | 3 | 622.909822 | 26.933619 | 0.796268 |

Table 8: Size test on NVIDIA GeForce RTX 2060 SUPER using 3 hash functions

| s | t | MOPS | Time (ms) | StdDev |
|------|-----|--------------------|-----------|----------|
| 2.0 | 2 | 4495.988473 | 3.731597 | 0.047658 |
| 1.9 | 2 | 4473.218103 | 3.750592 | 0.043783 |
| 1.8 | 2 | 4444.836303 | 3.774541 | 0.045024 |
| 1.7 | 2 | 2150.872191 | 7.800192 | 0.070234 |
| 1.6 | 2 | 4367.152402 | 3.841683 | 0.043648 |
| 1.5 | 2 | 4288.401832 | 3.912230 | 0.090330 |
| 1.4 | 2 | 4071.855181 | 4.120288 | 0.041169 |
| 1.3 | 2 | 3994.070087 | 4.200531 | 0.038393 |
| 1.2 | 2 | 3886.171126 | 4.317158 | 0.040137 |
| 1.1 | 2 | 3738.228285 | 4.488013 | 0.003802 |
| 1.05 | 2 | N/A | N/A | N/A |
| 1.02 | 2 | N/A | N/A | N/A |
| 1.01 | 2 | N/A | N/A | N/A |

Table 9: Size test on NVIDIA TESLA A100 using 2 hash functions

| s | t | MOPS | Time (ms) | StdDev |
|------|-----|--------------------|-----------|----------|
| 2.0 | 3 | 4624.068615 | 3.628237 | 0.044546 |
| 1.9 | 3 | 4615.325053 | 3.635110 | 0.045759 |
| 1.8 | 3 | 4595.526630 | 3.650771 | 0.045174 |
| 1.7 | 3 | 4583.313615 | 3.660499 | 0.048979 |
| 1.6 | 3 | 4556.363571 | 3.682150 | 0.045424 |
| 1.5 | 3 | 4539.911130 | 3.695494 | 0.052675 |
| 1.4 | 3 | 4518.867180 | 3.712704 | 0.049443 |
| 1.3 | 3 | 4483.699920 | 3.741824 | 0.043908 |
| 1.2 | 3 | 4452.022667 | 3.768448 | 0.038733 |
| 1.1 | 3 | 4416.581690 | 3.798688 | 0.038431 |
| 1.05 | 3 | 4381.700577 | 3.828928 | 0.044611 |
| 1.02 | 3 | 4364.069773 | 3.844397 | 0.044340 |
| 1.01 | 3 | 4328.415614 | 3.876064 | 0.053024 |

Table 10: Size test on NVIDIA TESLA A100 using 3 hash functions

duce reliable result, each experiment are repeated 5 times.

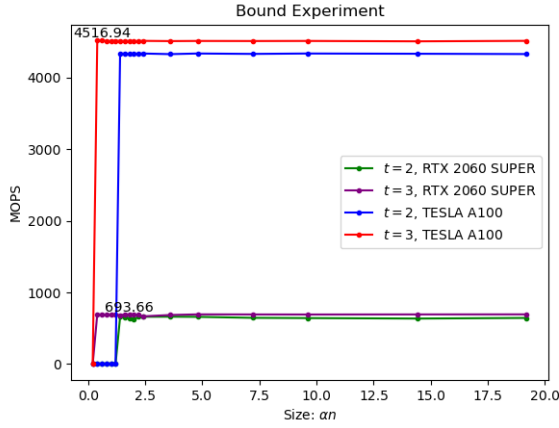


Figure 4: Experiment 4: Bound Test

| l | t | MOPS | Time (ms) | StdDev |
|------|-----|--------------------|-----------|----------|
| 0.2 | 2 | N/A | N/A | N/A |
| 0.4 | 2 | N/A | N/A | N/A |
| 0.6 | 2 | N/A | N/A | N/A |
| 0.8 | 2 | N/A | N/A | N/A |
| 1.0 | 2 | N/A | N/A | N/A |
| 1.2 | 2 | N/A | N/A | N/A |
| 1.4 | 2 | 4336.226983 | 3.869082 | 0.040889 |
| 1.6 | 2 | 4336.965930 | 3.868422 | 0.041448 |
| 1.8 | 2 | 4329.588092 | 3.875014 | 0.037551 |
| 2.0 | 2 | 4331.934797 | 3.872915 | 0.037947 |
| 2.2 | 2 | 4335.932864 | 3.869344 | 0.038864 |
| 2.4 | 2 | 4336.851103 | 3.868525 | 0.040728 |
| 3.6 | 2 | 4328.079694 | 3.876365 | 0.038323 |
| 4.8 | 2 | 4336.937169 | 3.868448 | 0.042185 |
| 7.2 | 2 | 4331.211905 | 3.873562 | 0.038510 |
| 9.6 | 2 | 4336.786582 | 3.868582 | 0.043544 |
| 14.4 | 2 | 4333.388402 | 3.871616 | 0.038374 |
| 19.2 | 2 | 4329.037487 | 3.875507 | 0.046952 |

Table 11: Bound test on NVIDIA TESLA A100 using 2 hash functions

From the Table 11, 12, 13, 14 and Figure 4, we can see that when choosing 2 hash functions, the performance remains almost the same after $l \geq 1.4$ (bound length $\geq 1.4 \log n$). The peak performances are 4336.965930 and 4516.142393 millions of operations per second, respectively.

When choosing 3 hash functions, we only need $l \geq 0.4$ to get a good performance. This may be because we choose a relative good hash function, xxHash, whose quality is excellent to easily pass Google hash test. And using 3 hash function itself makes the number of evictions much fewer than using 2 hash functions.

The best bound in our experiment is bolded in the table below, see Table 11, 12, 13, 14 and Figure 4 for more details.

| l | t | MOPS | Time (ms) | StdDev |
|------|-----|--------------------|-----------|----------|
| 0.2 | 3 | N/A | N/A | N/A |
| 0.4 | 3 | 4516.142393 | 3.714944 | 0.047456 |
| 0.6 | 3 | 4516.943821 | 3.714285 | 0.042790 |
| 0.8 | 3 | 4511.129706 | 3.719072 | 0.048300 |
| 1.0 | 3 | 4511.479023 | 3.718784 | 0.050594 |
| 1.2 | 3 | 4514.874533 | 3.715987 | 0.046709 |
| 1.4 | 3 | 4511.129590 | 3.719072 | 0.050060 |
| 1.6 | 3 | 4511.533343 | 3.718739 | 0.048303 |
| 1.8 | 3 | 4511.595532 | 3.718688 | 0.049938 |
| 2.0 | 3 | 4509.577610 | 3.720352 | 0.051675 |
| 2.2 | 3 | 4509.236274 | 3.720634 | 0.050588 |
| 2.4 | 3 | 4513.382040 | 3.717216 | 0.050051 |
| 3.6 | 3 | 4510.206033 | 3.719834 | 0.051105 |
| 4.8 | 3 | 4512.162369 | 3.718221 | 0.051948 |
| 7.2 | 3 | 4511.051972 | 3.719136 | 0.047064 |
| 9.6 | 3 | 4512.240083 | 3.718157 | 0.048978 |
| 14.4 | 3 | 4506.902747 | 3.722560 | 0.047575 |
| 19.2 | 3 | 4513.622961 | 3.717018 | 0.050403 |

Table 12: Bound test on NVIDIA TESLA A100 using 3 hash functions

| l | t | MOPS | Time (ms) | StdDev |
|------|-----|-------------------|-----------|----------|
| 0.2 | 2 | N/A | N/A | N/A |
| 0.4 | 2 | N/A | N/A | N/A |
| 0.6 | 2 | N/A | N/A | N/A |
| 0.8 | 2 | N/A | N/A | N/A |
| 1.0 | 2 | N/A | N/A | N/A |
| 1.2 | 2 | N/A | N/A | N/A |
| 1.4 | 2 | 657.483519 | 25.517318 | 0.312952 |
| 1.6 | 2 | 646.936886 | 25.933312 | 0.659120 |
| 1.8 | 2 | 641.007204 | 26.173210 | 0.573636 |
| 2.0 | 2 | 620.847396 | 27.023092 | 0.606418 |
| 2.2 | 2 | 658.149056 | 25.491514 | 0.317434 |
| 2.4 | 2 | 659.814903 | 25.427155 | 0.147894 |
| 3.6 | 2 | 661.132839 | 25.376468 | 0.109419 |
| 4.8 | 2 | 658.687197 | 25.470688 | 0.214283 |
| 7.2 | 2 | 644.927489 | 26.014112 | 0.493037 |
| 9.6 | 2 | 641.568838 | 26.150298 | 0.707538 |
| 14.4 | 2 | 634.861433 | 26.426579 | 0.982739 |
| 19.2 | 2 | 643.584807 | 26.068384 | 0.742324 |

Table 13: Bound test on NVIDIA 2060 SUPER using 2 hash functions

3.5 Comparison with Others

I compare my program's performance with some open source code found in github. From Figure at the top of result section, we can find that our program obtain a much higher performance both in insertion and lookup than all the open source codes I compared.

All the codes are tested in the same environment with GPU of NVIDIA RTX 2060 SUPER except for the first column (the result of my code on

| l | t | MOPS | Time (ms) | StdDev |
|------|-----|-------------------|-----------|----------|
| 0.2 | 3 | N/A | N/A | N/A |
| 0.4 | 3 | 693.425597 | 24.194688 | 0.115743 |
| 0.6 | 3 | 693.467047 | 24.193242 | 0.116118 |
| 0.8 | 3 | 690.305337 | 24.304051 | 0.234168 |
| 1.0 | 3 | 693.034023 | 24.208358 | 0.143322 |
| 1.2 | 3 | 693.265324 | 24.200282 | 0.120261 |
| 1.4 | 3 | 681.274671 | 24.626214 | 0.550600 |
| 1.6 | 3 | 693.066273 | 24.207232 | 0.109783 |
| 1.8 | 3 | 693.658439 | 24.186567 | 0.115933 |
| 2.0 | 3 | 691.121569 | 24.275347 | 0.161366 |
| 2.2 | 3 | 685.183113 | 24.485741 | 0.552498 |
| 2.4 | 3 | 665.937249 | 25.193389 | 0.589150 |
| 3.6 | 3 | 684.540419 | 24.508730 | 0.585207 |
| 4.8 | 3 | 693.442664 | 24.194093 | 0.113214 |
| 7.2 | 3 | 692.317996 | 24.233396 | 0.213677 |
| 9.6 | 3 | 691.614976 | 24.258029 | 0.131916 |
| 14.4 | 3 | 692.744288 | 24.218483 | 0.133526 |
| 19.2 | 3 | 693.520618 | 24.191373 | 0.115917 |

Table 14: Bound test on NVIDIA 2060 SUPER using 3 hash functions

NVIDIA TESLA A100).

References

- [1] Moulika Chadavada. Improving cuckoo hashing with perfect hashing. 2017.
- [2] Nathan Cassee and Anton Wijs. Analysing the performance of gpu hash tables for state space exploration. *Electronic Proceedings in Theoretical Computer Science*, 263:115, Dec 2017.
- [3] Salvatore Pontarelli, Pedro Reviriego, and Juan Antonio Maestro. Parallel d-pipeline: A cuckoo hashing implementation for increased throughput. *IEEE Transactions on Computers*, 65(1):326–331, 2016.
- [4] Yann Collet. xxhash - extremely fast non-cryptographic hash algorithm. 2016.